

# Introduction

This project explores apprenticeship contract cancellations in Australia using NCVER data. The aim is to:

1. Identify key drivers behind cancellations through EDA
2. Build a ML model to predict the likelihood of cancellation based on the available apprentice features

## 1. Prepare and clean data

```
In [1]: import pandas as pd
```

```
In [2]: #import tables
df_cancel = pd.read_excel("/Users/joelche/Documents/Projects/Apprenticeships/Apprentice_trainee.xlsx",sheet_nam
df_commence = pd.read_excel("/Users/joelche/Documents/Projects/Apprenticeships/Apprentice_trainee.xlsx",sheet_n
df_complete = pd.read_excel("/Users/joelche/Documents/Projects/Apprenticeships/Apprentice_trainee.xlsx",sheet_n
```

```
In [3]: #function to clean and reshape apprenticeship datasets

def clean_data(df_raw, status_type):

    #map 'type' to cancel flag
    cancel_flag = 1 if status_type == 'cancel' else 0

    #extract first 5 rows and transpose to get header columns
    multi_col_tuples = list(zip(
        df_raw.iloc[0,3:], #industry
        df_raw.iloc[1,3:], #occupation
        df_raw.iloc[2,3:], #gender
        df_raw.iloc[3,3:], #year
        df_raw.iloc[4,3:], #quarter
    ))

    #extract values
    #drop second column
    df_values = df_raw.drop(df_raw.columns[[0,2]],axis=1).reset_index(drop=True)
    #extract values for each state
    state_tuples = list(zip(
        df_values.iloc[5,1:], #totals
        df_values.iloc[6,1:], #NSW
        df_values.iloc[7,1:], #VIC
        df_values.iloc[8,1:], #QLD
        df_values.iloc[9,1:], #SA
        df_values.iloc[10,1:], #WA
        df_values.iloc[11,1:], #TAS
        df_values.iloc[12,1:], #NT
        df_values.iloc[13,1:], #ACT
    ))

    df_dim = pd.DataFrame(multi_col_tuples, columns=['industry','occupation','gender','year','quarter'])
    df_state = pd.DataFrame(state_tuples, columns=['totals','NSW','VIC','QLD','SA','WA','TAS','NT','ACT'])

    #concatenate dataframes
    df_raw = pd.concat([df_dim,df_state],axis=1)

    #add binary cancellation column where 1=cancelled and 0=not cancelled
    df_raw['cancel'] = cancel_flag

    #replace "-" with 0
    df_raw = df_raw.replace('-',0)

    #add
    df_raw['type'] = status_type

    return df_raw
```

```
In [4]: #apply function to raw data frames
df1_cancel = clean_data(df_cancel, 'cancel')
df1_commence = clean_data(df_commence, 'commence')
df1_complete = clean_data(df_complete, 'complete')
```

```
In [5]: #merge common fields
df_all = pd.concat([df1_cancel, df1_commence,df1_complete], ignore_index=True)
```

```
In [6]: #remove totals column
df_all = df_all.drop('totals',axis=1)
```

```
In [7]: #melt state columns into long format
```

```
df1_all = df_all.melt(
    id_vars=['occupation', 'gender', 'year', 'quarter', 'cancel', 'type'],
    value_vars=['NSW', 'VIC', 'QLD', 'SA', 'WA', 'TAS', 'NT', 'ACT'],
    var_name='state',
    value_name='count'
)
```

```
In [8]: #since we are removing the industry column we have to group and aggregate values
df2_all = df1_all.groupby(['occupation', 'gender', 'year', 'quarter', 'state', 'cancel', 'type'], as_index=False)['co
```

```
In [9]: #remove all rows where count==0
df2_all = df2_all[df2_all['count'] > 0]
```

```
In [10]: df2_all.head(10)
```

```
Out[10]:
```

	occupation	gender	year	quarter	state	cancel	type	count
0	1112 - General Managers	Females	2019	Apr-Jun	ACT	0	complete	5
5	1112 - General Managers	Females	2019	Apr-Jun	TAS	0	complete	5
7	1112 - General Managers	Females	2019	Apr-Jun	WA	0	complete	5
8	1112 - General Managers	Females	2019	Jul-Sep	ACT	0	complete	10
14	1112 - General Managers	Females	2019	Jul-Sep	VIC	0	complete	5
26	1112 - General Managers	Females	2019	Oct-Dec	TAS	0	commence	10
44	1112 - General Managers	Females	2020	Apr-Jun	SA	0	commence	5
54	1112 - General Managers	Females	2020	Apr-Jun	WA	0	complete	5
64	1112 - General Managers	Females	2020	Jul-Sep	ACT	0	complete	10
65	1112 - General Managers	Females	2020	Jul-Sep	ACT	1	cancel	5

```
In [11]: #check for duplicate values
dups = df2_all.duplicated(subset=['occupation', 'gender', 'year', 'quarter', 'state', 'cancel'])
print("Remaining duplicates:", dups.sum())
```

Remaining duplicates: 4879

```
In [12]: df2_all.shape
```

```
Out[12]: (28584, 8)
```

```
In [13]: df2_all.nunique()
```

```
Out[13]:
```

occupation	212
gender	3
year	6
quarter	4
state	8
cancel	2
type	3
count	198
dtype:	int64

```
In [14]: #remove unknown gender
df2_all = df2_all[df2_all['gender'].isin(['Males', 'Females'])]
```

```
In [15]: df2_all['gender'].unique()
```

```
Out[15]: array(['Females', 'Males'], dtype=object)
```

## 2. Exploratory Data Analysis

In the exploratory data analysis, I will analyse each feature individually and make an assessment on the impact it has on the cancellation rate of apprenticeship programs

The cancellation rate can be defined as the percentage of apprentices who cancelled their training contracts out of all apprentices who are no longer in training (i.e. either cancelled or completed)

### 2.1 Gender analysis

```
In [16]: import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [17]: #cancellation rate for each gender
df_end = df2_all[df2_all['type'].isin(['cancel', 'complete'])]

#group and sum counts by gender and type
```

```
gender_summary = df_end.groupby(['gender', 'type'])['count'].sum().unstack(fill_value=0)

#calculate cancellation rate
gender_summary['cancel_rate'] = gender_summary['cancel']/(gender_summary['cancel'] + gender_summary['complete'])

gender_summary
```

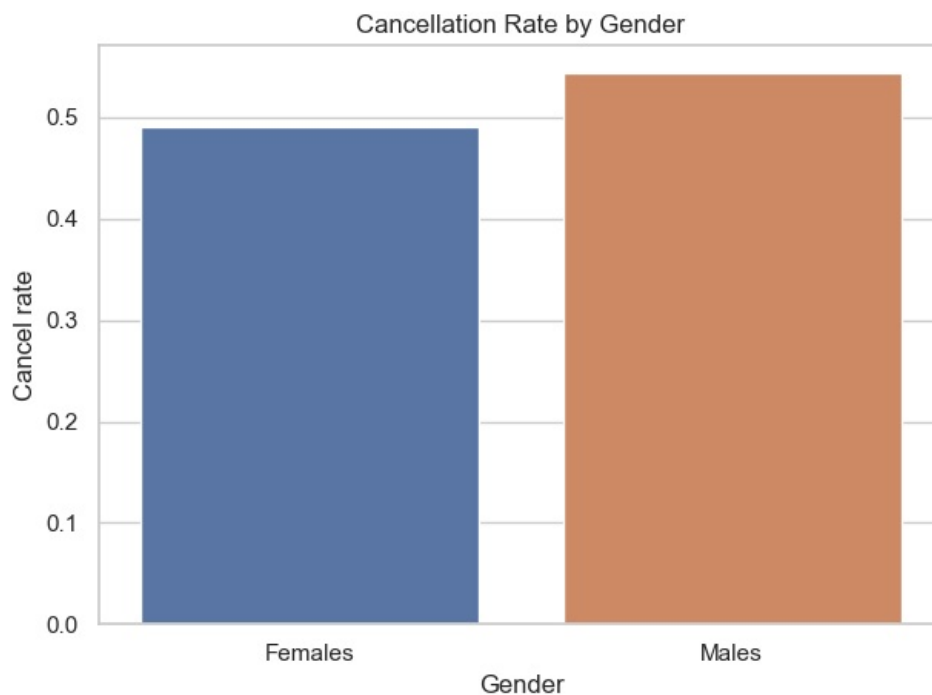
```
Out[17]:
```

	type	cancel	complete	cancel_rate
gender				
Females		75255	78235	0.490293
Males		249080	209155	0.543564

```
In [18]: #plot results
sns.set(style='whitegrid')
ax = sns.barplot(x=gender_summary.index, y=gender_summary['cancel_rate'])

#add labels
ax.set_title('Cancellation Rate by Gender')
ax.set_ylabel('Cancel rate')
ax.set_xlabel('Gender')

plt.tight_layout()
plt.show()
```



Key insights:

1. Males(54%) have slightly higher cancellation rate than females(49%) suggesting males are more likely to not complete their apprenticeships than females
2. Both genders have high cancellation rates. This could indicate structural issues with support, expectations, job alignment etc.
3. There are far more males than females participating in apprenticeship programs. This reflects that males still make up the majority of trade employment

## 2.2 Occupation analysis

```
In [19]: #group and sum counts by occupation and type
occ_summary = df_end.groupby(['occupation', 'type'])['count'].sum().unstack(fill_value=0)

#calculate cancellation rate
occ_summary['cancel_rate'] = occ_summary['cancel']/(occ_summary['cancel'] + occ_summary['complete'])

#absolute count
occ_summary['abs_count'] = occ_summary['cancel']+occ_summary['complete']

#sort by highest cancellation rate
occ_summary = occ_summary.sort_values(by='cancel_rate', ascending=False)

#exclude low-volume occupations
occ_summary = occ_summary[(occ_summary['cancel']+occ_summary['complete'])>1000]

occ_summary.head(20)
```

Out[19]:

	type	cancel	complete	cancel_rate	abs_count
occupation					
2721 - Counsellors		1725	440	0.796767	2165
5911 - Purchasing and Supply Logistics Clerks		7645	2935	0.722590	10580
4314 - Hotel Service Managers		2150	880	0.709571	3030
3126 - Safety Inspectors		1015	455	0.690476	1470
3121 - Architectural, Building and Surveying Technicians		1750	800	0.686275	2550
3511 - Bakers and Pastrycooks		2810	1405	0.666667	4215
3333 - Roof Tilers		985	495	0.665541	1480
3514 - Cooks		12420	6265	0.664704	18685
3321 - Floor Finishers		695	355	0.661905	1050
6214 - Pharmacy Sales Assistants		2915	1570	0.649944	4485
3627 - Landscape Gardeners and Irrigation Technicians		4785	2595	0.648374	7380
6215 - Retail Supervisors		1060	580	0.646341	1640
6121 - Real Estate Sales Agents		8190	4630	0.638846	12820
8219 - Other Construction and Mining Labourers		2075	1215	0.630699	3290
3311 - Bricklayers and Stonemasons		3215	1900	0.628543	5115
5121 - Office Managers		7175	4260	0.627460	11435
3322 - Painters		4015	2500	0.616270	6515
4310 - Hospitality Workers - nfd		24200	15100	0.615776	39300
3332 - Plasterers and Renderers		2055	1300	0.612519	3355
3341 - Plumbers		28055	18540	0.602103	46595


In [20]: #plot results

```
top_occ = occ_summary[occ_summary['cancel_rate']>0.5].head(20)

ax = sns.barplot(data=top_occ.reset_index(),y='occupation',x='cancel_rate',palette='Reds_r')

#add labels
ax.set_title('Top 20 Occs by Cancel Rate')
ax.set_ylabel('Occupation')
ax.set_xlabel('Cancel rate')

plt.xlim(0,1)
plt.tight_layout()
plt.show()
```



Key insights:

1. Hospitality and retail are overrepresented. Cooks, retail supervisors, pharmacy sales assistant all appear in the top 20. This can be due to the industry having high staff turnover, involve shift work, high labour and low pay.

- Construction trades such as roof tilers and painters also appear heavily. Possibly driven by physically demanding work or poor work culture?
- High volume count occupations such as plumbers and cooks should be priority targets for rententio strategies

## 2.3 Location analysis

```
In [21]: state_summary = df_end.groupby(['state', 'type'])['count'].sum().unstack(fill_value=0)

# Calculate cancellation rate
state_summary['cancel_rate'] = state_summary['cancel'] / (state_summary['cancel'] + state_summary['complete'])

#absolute count
state_summary['abs_count'] = state_summary['cancel']+state_summary['complete']

# Sort by cancellation rate (optional)
state_summary = state_summary.sort_values('cancel_rate', ascending=False)

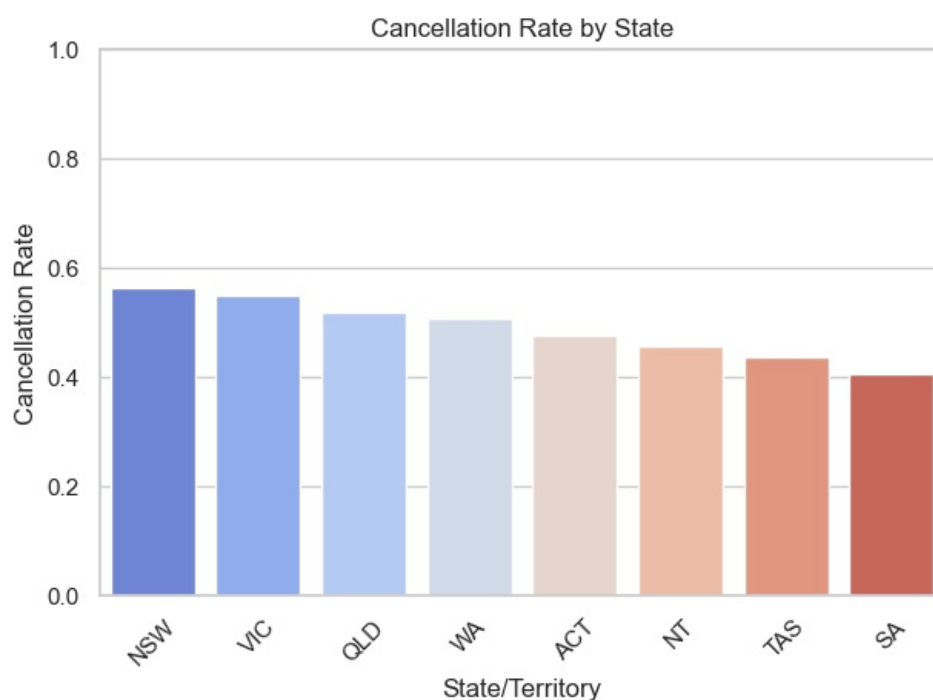
state_summary
```

```
Out[21]:
```

	type	cancel	complete	cancel_rate	abs_count
state					
NSW		103415	79650	0.564909	183065
VIC		81790	67085	0.549387	148875
QLD		79680	73475	0.520257	153155
WA		32645	31735	0.507067	64380
ACT		5035	5500	0.477931	10535
NT		2050	2445	0.456062	4495
TAS		7040	9075	0.436860	16115
SA		12680	18425	0.407652	31105

```
In [22]: #plot results
# Reset index for plotting
state_plot = state_summary.reset_index()

sns.barplot(data=state_plot, x='state', y='cancel_rate', palette='coolwarm')
plt.title('Cancellation Rate by State')
plt.ylabel('Cancellation Rate')
plt.xlabel('State/Territory')
plt.ylim(0, 1)
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```



Key insights:

- Larger states (NSW, VIC, QLD) have higher cancellation rates which suggest systematic challenges in managing apprenticeship support and training quality in areas with larger population

2. Smaller states have lower cancel rate which could be due to more manageable apprentice-to-support ratios and less occupational variety

### 3. Build ML model

I will build a machine learning model to better understand apprenticeship cancellation in Australia. The model will help identify key drivers of cancellation - such as occupation, gender, and location - and predict which apprentices are most at risk of cancelling.

```
In [23]: #import libraries
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report, confusion_matrix, roc_auc_score
from sklearn.linear_model import LogisticRegression
```

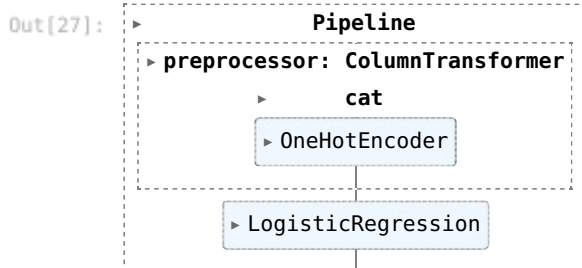
```
In [24]: #Define features and target
categorical_features = ['occupation', 'gender', 'year', 'quarter', 'state']
X = df2_all[categorical_features]
y = df2_all['cancel']
weights = df2_all['count'] # Need pipeline to account for the count of each row value
```

```
In [25]: #Preprocessing (OneHotEncoder for categorical variables)
preprocessor = ColumnTransformer(
    transformers=[
        ('cat', OneHotEncoder(handle_unknown='ignore'), categorical_features)
    ],
    remainder='drop' # No numeric features to passthrough
)
```

```
In [26]: # Create pipeline
model = Pipeline(steps=[
    ('preprocessor', preprocessor),
    ('classifier', LogisticRegression(max_iter=1000, class_weight=None)) # We'll use sample_weight instead
])
```

```
In [27]: # Train-test split
X_train, X_test, y_train, y_test, weights_train, weights_test = train_test_split(
    X, y, weights, test_size=0.2, stratify=y, random_state=42
)

# Fit model with weights
model.fit(X_train, y_train, classifier__sample_weight=weights_train)
```




```
In [28]: # Evaluate model
y_pred = model.predict(X_test)
y_proba = model.predict_proba(X_test)[: , 1]

print(classification_report(y_test, y_pred))
print("ROC AUC:", roc_auc_score(y_test, y_proba))

# Plot confusion matrix
conf_matrix = confusion_matrix(y_test, y_pred)
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues')
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix")
plt.show()
```

	precision	recall	f1-score	support
0	0.68	0.93	0.79	3776
1	0.53	0.16	0.25	1941
accuracy			0.67	5717
macro avg	0.60	0.54	0.52	5717
weighted avg	0.63	0.67	0.60	5717

ROC AUC: 0.6158241891083576



Summary of ML model:

We trained the logistic regression model to predict the likelihood of an apprentice cancelling their training contract based on the available features. The model performs well in predicting completions (68% precision) however struggles to identify cancellations (53% precision). ROC AUC score is 0.62 suggesting model has limited ability to distinguish between apprentice who cancel and those who complete. This may be due to insufficient predictive features.

## Conclusion

### EDA Findings

Gender: Males had slightly higher cancellation rate (54%) compared to females (49%) State: NSW and VIC had the highest cancellation rates (~55-56%), while SA had the lowest (41%) Occupation: Roles in the hospitality sector had notably high cancellation rates

### ML Model

The model struggled to identify cancelled apprenticeships (many false negatives). Suggest re testing with different model or using other predictive features.